

Hidden Layers Used in Deep Learning

Introduction :

In deep learning, a hidden layer is a layer between the input layer and output layer. These layers perform complex computations and help the model learn patterns, features, and representations from data. A neural network with more than one hidden layer is called a Deep Neural Network (DNN).

Consists of three primary types of layers:

- Input Layer
- Hidden Layers
- Output Layer

Types of Hidden Layers :

1. Dense (Fully Connected) Layer :

Dense (Fully Connected) Layer is the most common type of hidden layer in an ANN. Every neuron in a dense layer is connected to every neuron in the previous and subsequent layers.

- Role: Learns representations from input data.
- Function: Performs weighted sum and activation.

2. Convolutional Layer

Convolutional layers are used in Convolutional Neural Networks (CNNs) for image processing tasks. They apply convolution operations to the input, capturing spatial hierarchies in the data.

- Role: Extracts spatial features from images.
- Function: Applies convolution using filters

3. Recurrent Layer

Recurrent layers are used in Recurrent Neural Networks (RNNs) for sequence data like time series or natural Language.

- Role: Processes sequential data with temporal dependencies.
- Function: Maintains state across time steps.

4. Dropout Layer

Dropout layers are a regularization technique used to prevent overfitting.

- Role: Prevents overfitting.
- Function: Randomly drops neurons during training.

5. Pooling Layer

Pooling Layer is used to reduce the spatial dimensions of the data, thereby decreasing the computational load and controlling overfitting.

- Use Cases: Dimensionality reduction in CNNs

6. Batch Normalization Layer

A Batch Normalization Layer normalizes the output of a previous activation layer by subtracting the batch mean and dividing by the batch standard deviation.

- Use Cases: Stabilizing and speeding up training

Activation Functions Used in the Deep learning

1. Linear Activation Function

Formula:

- Output is same as input
- Used in regression output layer

- No non-linearity

1. Binary Step Function

- Used in early Perceptron
- Not used in modern deep learning (not differentiable)

2. Sigmoid (Logistic) Function

- Output range: 0 to 1
- Used in binary classification (output layer)
- Problem: Vanishing gradient

4. Tanh (Hyperbolic Tangent)

- Output range: -1 to +1
- Zero-centered
- Still suffers from vanishing gradient

5. ReLU (Rectified Linear Unit)

- Most widely used
- Faster training
- Problem: Dead neuron problem

6. Leaky ReLU

- Solves dead neuron problem
- Small slope for negative values

7. Parametric ReLU (PReLU)

- A is learnable parameter
- Improved version of Leaky ReLU

8. ELU (Exponential Linear Unit)

- Smooth curve for negative values
- Reduces vanishing gradient
- Faster convergence than ReLU

9. SELU (Scaled ELU)

- Self-normalizing property
- Used in deep networks

10. Softmax

- Used in multi-class classification output layer
- Output range: 0 to 1 (probabilities)

To Find metrics, optimiser, losses

1. Loss Functions (Objectives to Minimize)

Regression Tasks:

- Mean Squared Error (MSE/L2 Loss): Used when outliers are not a major concern.
- Mean Absolute Error (MAE/L1 Loss): More robust to outliers.

Classification Tasks:

- Binary Cross-Entropy (Log Loss): For binary classification.
- Categorical Cross-Entropy: For multi-class classification.
- Sparse Categorical Cross-Entropy: Used when labels are integers rather than one-hot encoded.
- Hinge Loss: Used frequently in Support Vector Machines (SVMs).

2.Optimizers (Algorithms to Update Weights)

- Adam (Adaptive Moment Estimation): A popular, generally effective optimizer that combines the strengths of AdaGrad and RMSProp.
- SGD (Stochastic Gradient Descent): Updates parameters in the opposite direction of the gradient, often used with momentum.

- AdamW: A variant of Adam that implements decoupled weight decay, often used to improve generalization in Transformer models.
- RMSProp: Adaptive algorithm for handling sparse/noisy datasets.
- Adafactor: Memory-efficient optimizer used for very large-scale models.

3. Metrics (Evaluation Criteria)

Classification:

- Accuracy: Overall accuracy or sparse categorical accuracy for integer labels.
- Precision, Recall, F1 Score: Essential for imbalanced datasets.
- AUC-ROC: Measures performance across different thresholds.
- Confusion Matrix: Provides a detailed breakdown of correct/incorrect predictions.

Regression:

- MAE (Mean Absolute Error)
- MSE (Mean Squared Error)
- RMSE (Root Mean Squared Error)

NLP/Sequence:

- BLEU Score
- ROUGE Score
- Perplexity

